

COMP3605 Introduction to Data Analytics Cheat Sheet

based on lecture notes by Dr. Duc The Kieu

Christopher Sahadeo

2018-10-13

Contents

I	Key Terms	3
II	Algorithms, Formulae and Examples	6
1	Measures	6
1.1	Distance Metrics	6
1.2	Means	7
1.3	Variance	7
2	Classification	8
2.1	Decision Trees	8
2.1.1	Information Gain (ID3)	9
2.1.2	Gain Ratio (C4.5)	9
2.1.3	Gini Index (CART)	9
2.2	Naïve Bayesian Classification	10
2.3	Rule-Based Classification TBC	11
2.4	Evaluate Classifier Performance	12
3	Advanced Classification	13
3.1	k-Nearest Neighbors Classification	13
3.2	Classification Using Frequent Patterns	13
3.3	Support Vector Machines (SVMs)	13

3.4	Classification by Backpropagation (ANNs)	13
3.5	Bayesian Belief Networks	13
3.6	Other Classification Methods	13
4	Clustering	14
4.1	k-Means Algorithm (partitioning based)	14
4.2	DBSCAN Algorithm (density based) TBC	15
4.3	Evaluate Clustering Performance	16
5	Outlier Detection	17
5.1	Classification-Based Approaches	17
5.2	Clustering-Based Approaches	17
5.3	Proximity-Based Approaches	17
5.4	Distance-Based Approaches	17
5.5	Density-Based Approaches	17
5.6	Statistical Approaches	17

Part I

Key Terms

Data classification	A two-step process, consisting of a learning (or training) step (where a classification model is constructed) and a classification step (where the model is used to predict class labels for given data)
Supervised Learning	The class label of each training tuple is provided
Unsupervised Learning	The class label of each training tuple is not known
Accuracy	the percentage of test set tuples that are correctly classified by the (trained/learned) classifier
D	(parent set) is training set of class-labeled tuples
$X_{train} = (x_1, \dots, x_n, y)$	A tuple in D ; an attribute/feature
y	class label attribute
D_{Test}	Test data containing tuples $X_{test} = (x_1, \dots, x_n)$ without the class attribute
m	The number of classes in D , each class is denoted as C_i (for $i = 1, \dots, m$)
$C_{i,D}$	The set of tuples of class C_i in D
$ D $	The number of tuples in D
$ C_{i,D} $	The number of tuples in $C_{i,D}$
v	The number of distinct values of attribute A
a_j	A given value of attribute A (for $j = 1, \dots, v$)
D_j	The set of tuples in D that have outcome a_j of A
p_i	The nonzero probability that an arbitrary tuple in D belongs to class C_i
$Info(D)$	Expected information needed to identify the class label of a tuple, before partitioning on A
$Info(D_j)$	Expected information needed to identify the class label of a tuple, after partitioning on A
$Info_A(D)$	Actual information still needed to identify the class label of a tuple, after partitioning on A
$Gain(A)$	The attribute with the highest value is chosen as the splitting attribute, biased toward tests with many outcomes
$GainRatio(A)$	The attribute with the highest value is chosen as the splitting attribute, overcomes the Information Gain bias, but it prefers unbalanced splits in which one partition is much smaller than the others
$Gini(D)$	The subset D_1 or D_2 , upon binary split, that gives minimum Gini index for that attribute is selected as its splitting subset, overcomes the Gain Ratio bias, but is biased to multivalued attributes
H	A hypothesis that tuple X belongs to a class C
$P(H X)$	posterior probability, posteriori probability the probability that tuple X belongs to class C , given that we know the attribute description of X
$P(X H)$	The posterior probability that we can determine the description of X given that we know X belongs to class C

$P(H)$	prior probability, priori probability the probability that X belongs to C regardless of the description of X
$P(X)$	A constant the prior probability that we can determine the description of X regardless of what class X belongs to
$P(x_k C_i)$	The number of tuples of class C_i in D having the value x_k for categorical attribute A_k , divided by $ C_{i,D} $, the number of tuples of class C_i in D
R	A Rule, R : IF <i>condition</i> THEN <i>conclusion</i>
n_{covers}	The number of tuples covered by R , if the condition is satisfied
$n_{correct}$	The number of tuples correctly classified by R
$coverage(R)$	The percentage of tuples that are covered by the rule (i.e., their attribute values hold true for the rule antecedent)
$accuracy(R)$	The percentage of tuples (covered by the rule) that are correctly classified
TP	The number of positive tuples that were correctly labeled by the classifier
TN	The number of negative tuples that were correctly labeled by the classifier
FP	The number of negative tuples that were mislabeled as positive
FN	The number of positive tuples that were mislabeled as negative
Rule-Based Classification	Learned/trained model is represented as set of IF-THEN rules. Uses either decision tree induction or sequential covering algorithm .
Rule coverage (satisfied)	If condition (i.e., all the attribute tests) in rule antecedent holds true for given tuple. If a rule is satisfied, it is said to be triggered
Conflict Resolution	Two or more rules are triggered, we need a conflict resolution strategy to figure out which rule gets to fire and assign its class prediction to X
Size ordering	Assigns highest priority to triggering rule that has the toughest requirements
Classed-based Rule ordering	All the rules for the most prevalent (or most frequent) classes come first
Rule-based Rule ordering	Rules are organized into one long priority list, according to some measure of rule quality
Clustering	The process of grouping set of data objects into multiple clusters (or groups) so that objects within a cluster have high similarity, but are very dissimilar to objects in other clusters
cluster	A collection of data objects that are similar to one another within the cluster and dissimilar to objects in other clusters
p	An object in a cluster
c_i	The centroid (center/mean) of a cluster
m_i	The number of objects in cluster i
$dist(x, y)$	The Euclidean Distance between two points $x = (x_1, \dots, x_d)$ and $y = (y_1, \dots, y_d)$
E	Sum of squared error (SSE) between all objects in C_i and the centroid c_i
k -Means Algorithm	Distributes the objects in D into a set of k clusters C_1, \dots, C_k such that $C_i \subset D$ and $C_i \cap C_j = \emptyset$ for $1 \leq i, j \leq k$, designed to find spherical shaped clusters
DBSCAN	Density-Based Spatial Clustering of Applications with Noise, finds clusters of arbitrary shape, we can model clusters as dense regions in the data space, separated by sparse regions , discovers clusters of non-spherical shape, finds core objects (i.e., objects that have dense neighborhoods)
density	Number of objects close to o

ϵ -neighborhood	The space (region, area) within a radius ϵ centered at o , where $\epsilon > 0$ is a user-specified parameter
N	The set of objects in the ϵ -neighborhood of a point p under consideration
$ N $	The number of elements in N
core object (point)	p is a core point if $ N \geq MinPts$
Border point	q is not a core point (i.e., its ϵ -neighborhood contains less than $MinPts$ points, $ N < MinPts$), but falls within the ϵ -neighborhood of a core point (i.e., its ϵ -neighborhood contains at least one core point)
Noise point	Any point that is neither a core point nor a border point

Part II

Algorithms, Formulae and Examples

1 Measures

1.1 Distance Metrics

Minkowski Denoted as $L_p(x, y), L^p(x, y), p \geq 2$

$$d_{Minkowski}(x, y) = \left(\sum_{j=1}^d |y_j - x_j|^p \right)^{\frac{1}{p}}$$

Manhattan (Minkowski with $p = 1$), denoted as $L_1(x, y), L^1(x, y)$

$$d_{Manhattan}(x, y) = \sum_{j=1}^d |y_j - x_j|$$

Euclidean (Minkowski with $p = 2$), denoted as $L_2(x, y), L^2(x, y)$

$$d_{Euclidean}(x, y) = \sqrt{\sum_{j=1}^d (y_j - x_j)^2}$$

Cosine

$$\begin{aligned} x \cdot y &= \sum_{j=1}^d x_j y_j \\ \|x\| &= \sqrt{\sum_{j=1}^d x_j^2} \\ d_{cosine}(x, y) &= \frac{x \cdot y}{\|x\| \times \|y\|} \end{aligned}$$

Mahalanobis

$$d_{Mahalanobis}(x, y) = \sqrt{(x - y)^T \Sigma^{-1} (x - y)}$$

where

- Σ is a covariance matrix
- Σ^{-1} is the inverse of Σ
- x^T is the transpose of x

1.2 Means

1.3 Variance

2 Classification

2.1 Decision Trees

$$O(n \times |D| \times \log(|D|))$$

Algorithm 1 Generate_decision_tree

Input: D , $attribute_list$, $Attribute_selection_method \in \{information_gain, gain_ratio, gini_index\}$

Output: A decision tree

```
1: procedure GENERATE_DECISION_TREE( $D, attribute\_list, Attribute\_selection\_method$ )
2:   create a node  $N$ 
3:   if  $X \in C, \forall X \in D$  then
4:     return  $N$  as leaf node labeled with class  $C$ 
5:   end if
6:   if  $attribute\_list = \emptyset$  then
7:     return  $N$  as leaf node labeled with the majority class in  $D$ 
8:   end if
9:   apply ATTRIBUTE_SELECTION_METHOD( $D, attribute\_list$ ) to find the "best"  $splitting\_criterion$ 
10:  if  $splitting\_attribute$  is discrete and multiway splits allowed then
11:     $attribute\_list \leftarrow attribute\_list - splitting\_attribute$ 
12:  end if
13:  for each outcome  $j$  of  $splitting\_criterion$  do
14:    if  $D_j = \emptyset$  then
15:      attach a leaf labeled with the majority class in  $D$  to node  $N$ 
16:    else
17:      attach the node returned by GENERATE_DECISION_TREE( $D_j, attribute\_list$ ) to node  $N$ 
18:    end if
19:  end for
20:  return  $N$ 
21: end procedure
```

2.1.1 Information Gain (ID3)

$$\begin{aligned}p_i &= \frac{|C_{i,D}|}{|D|} \\Info(D) &= -\sum_{i=1}^m p_i \log_2(p_i) \\p_{i,j} &= \frac{|C_{i,D_j}|}{|D_j|} \\Info(D_j) &= -\sum_{i=1}^m p_{i,j} \log_2(p_{i,j}) \\Info_A(D) &= \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j) \\Gain(A) &= Info(D) - Info_A(D) \\0 \log_2 0 &= 0\end{aligned}$$

[Click here for Example](#)

2.1.2 Gain Ratio (C4.5)

$$\begin{aligned}SplitInfo_A(D) &= -\sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2\left(\frac{|D_j|}{|D|}\right) \\GainRatio(A) &= \frac{Gain(A)}{SplitInfo_A(D)}\end{aligned}$$

[Click here for Example](#)

2.1.3 Gini Index (CART)

$$Gini(D) = 1 - \sum_{i=1}^m p_i^2$$

Determines the splitting attribute and splitting subsets:

$$Gini_A(D) = \frac{|D_1|}{|D|} Gini(D_1) + \frac{|D_2|}{|D|} Gini(D_2)$$

Determines the reduction in impurity:

$$\Delta Gini(A) = Gini(D) - Gini_A(D)$$

[Click here for Example](#)

2.2 Naïve Bayesian Classification

$$P(x_k|C_i) = \frac{|C_{i,x_k}|}{|C_{i,D}|}$$

Laplacian Correction if $P(x_k|C_i) = 0$: If we have q counts to which we each add one, then we must remember to add q to the corresponding denominator used in the probability calculation

$$P(X|C_i) = \prod_{k=1}^n P(x_k|C_i)$$

$$P(C_i|X) = \frac{P(X|C_i)P(C_i)}{P(X)}$$

Case 1 If $P(C_i)$ is known, then we calculate it by $P(C_i) = \frac{|C_{i,D}|}{|D|}$ and we maximize

$$P(X|C_i)P(C_i) > P(X|C_j)P(C_j) \text{ for } 1 \leq j \leq m, j \neq i$$

Case 2 If $P(C_i)$ is unknown, then we assume $P(C_i) = P(C_j)$ and we maximize

$$P(X|C_i) > P(X|C_j)$$

[Click here for Example](#)

2.3 Rule-Based Classification TBC

$$\text{coverage}(R) = \frac{n_{\text{covers}}}{|D|}$$
$$\text{accuracy}(R) = \frac{n_{\text{correct}}}{n_{\text{covers}}}$$

Algorithm 2 Sequential Covering Algorithm

Input: D, Att_vals

Output: A set of IF-THEN rules

```
1: procedure SEQUENTIAL COVERING( $D, Att\_vals$ )
2:    $R\_set = \{\}$ 
3:   for each class  $c$  do
4:     repeat
5:        $R = \text{LEARN\_ONE\_RULE}(D, Att\_vals, c)$ 
6:       remove tuples covered by  $R$  from  $D$ 
7:        $R\_set = R\_set + R$ 
8:     until terminating condition
9:   end for
10:  return  $R\_set$ 
11: end procedure
```

$$FOIL_Gain = pos' \times \left(\log_2 \frac{pos'}{pos' + neg'} - \log_2 \frac{pos}{pos + neg} \right)$$

2.4 Evaluate Classifier Performance

		Predicted Class		
		yes	no	Total
Actual class—	yes	TP	FN	P
	no	FP	TN	N
		Total	P'	N'
				P + N

Confusion Matrix

$$accuracy = \frac{TP + TN}{P + N}$$

$$error_rate = \frac{FP + FN}{P + N}$$

$$sensitivity, true_positive_rate, recall = \frac{TP}{P}$$

$$specificity, true_negative_rate = \frac{TN}{N}$$

$$precision = \frac{TP}{TP + FP}$$

$$F_1 = \frac{2 \times precision \times sensitivity}{precision + sensitivity}$$

$$F_\beta = \frac{(1 + \beta^2) \times precision \times sensitivity}{\beta^2 \times precision + sensitivity}, \beta \in \mathbb{R}, \beta \geq 0$$

[Click here for Example](#)

3 Advanced Classification

3.1 k-Nearest Neighbors Classification

3.2 Classification Using Frequent Patterns

3.3 Support Vector Machines (SVMs)

3.4 Classification by Backpropagation (ANNs)

3.5 Bayesian Belief Networks

3.6 Other Classification Methods

4 Clustering

4.1 k-Means Algorithm (partitioning based)

$$c_i = \frac{1}{m_i} \sum_{p \in C_i} p$$
$$dist(x, y) = \sqrt{\sum_{j=1}^d (y_j - x_j)^2}$$
$$E = \sum_{i=1}^k \sum_{p \in C_i} dist(p, c_i)^2$$

Algorithm 3 k-Means

Input: k, D

Output: A set of k clusters C_1, \dots, C_k such that $C_i \subset D$ and $C_i \cap C_j = \emptyset$ for $1 \leq i, j \leq k$

- 1: **procedure** $k\text{-MEANS}(k, D)$
 - 2: Randomly choose k objects from D as initial centroids
 - 3: **repeat**
 - 4: (re)assign each object to the cluster to which the object is the most similar
 - 5: recalculate the centroids
 - 6: **until** centroids do not change
 - 7: **end procedure**
-

[Click here for Example](#)

4.2 DBSCAN Algorithm (density based) TBC

Density-Based Spatial Clustering of Applications with Noise

4.3 Evaluate Clustering Performance

5 Outlier Detection

5.1 Classification-Based Approaches

5.2 Clustering-Based Approaches

5.3 Proximity-Based Approaches

5.4 Distance-Based Approaches

5.5 Density-Based Approaches

5.6 Statistical Approaches